PERSONAL AI SYSTEM

# Your personal AI workforce.

One human director. A coordinating AI. Specialist workers. Decisions that remain yours.

SUITE FLEET IN ONE SENTENCE

**You decide what matters. Muse coordinates the work. Specialized AI seats and tools do the work. You approve the decisions that count.**

## THE FOUR EVERYDAY JOBS

### 01 Email management

Sort and summarize incoming messages; identify what needs attention.

### 02 Job lead scoring

Rank job postings against salary, location, title and employer rules.

### 03 Deal evaluation

Check whether a deal is credible, relevant and worth acting on.

### 04 Notes quality checks

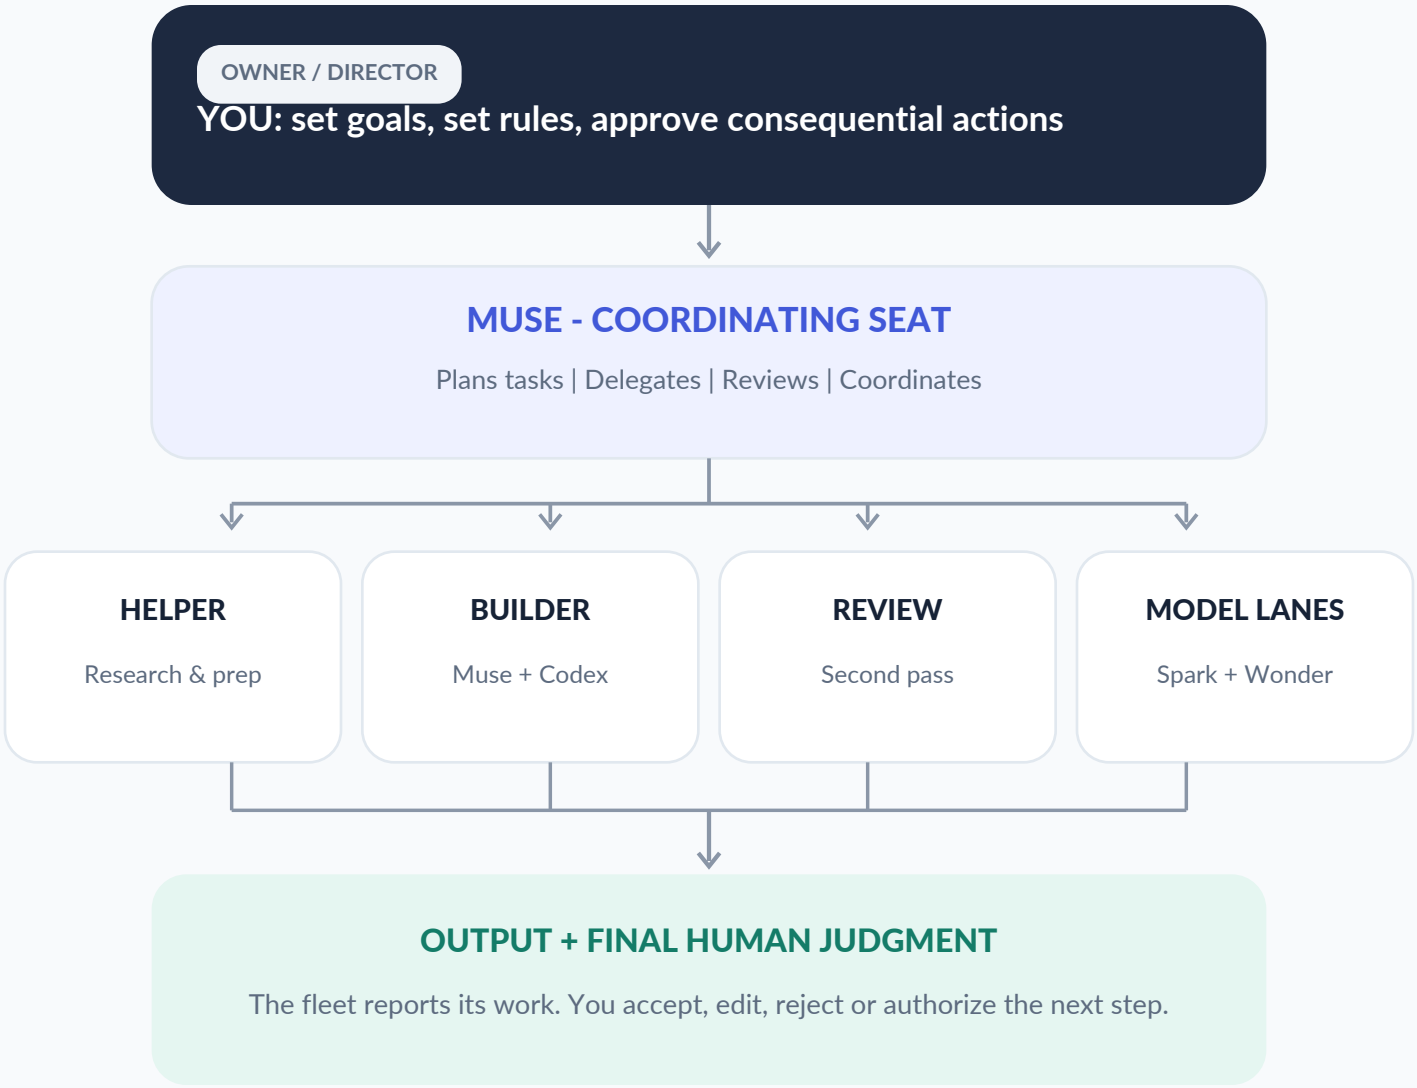
Look for stale facts, gaps and mistakes in your notes.

These are repeatable workflows, not just one-off questions sent to a chatbot.

02 / THE ORGANIZATION

# How the AI seats work together

Think of the system as a tiny organization - but you are always the final authority.




HARD APPROVAL GATES

**Spending money | Sending outward messages | Making submissions**

A seat is a responsibility in the system - not necessarily a different model talking all the time.

03 / BUILD FOR YOUR WORLD

# What the Meta hackathon added

The fleet already existed. The hackathon work is a routing-and-measurement layer for the same real tasks.

EXISTING FLEET → **NEW META MODEL ROUTING + BENCHMARKS**

Measures real work instead of showcasing only synthetic examples.

## SAME WORK, TWO CANDIDATE MODELS

### S MUSE SPARK 1.3

Evaluated on the same mail and job examples as Wonder.

48 logged calls

### W MUSE WONDER 1.0

Evaluated head-to-head, with its outputs preserved for review.

47 logged calls

## HOW EACH BENCHMARK DECISION BECOMES EVIDENCE

1. **REAL INPUT** — Email or job lead
2. **TWO OUTPUTS** — Spark and Wonder
3. **LEDGER** — Tokens, speed, cost
4. **REVIEW** — Accept, edit, reject

### What this lets you answer

Which model is faster? Which is cheaper under estimated pricing? How often do they agree? Which outputs would the owner actually accept?

Published benchmarks cover mail triage and job scoring. Deal evaluation and notes QA have no published measured runs yet.

04 / PUBLISHED RESULTS

# The numbers - and what they mean

Recorded benchmark activity from October 26-28, 2026, with owner verdicts applied October 30.

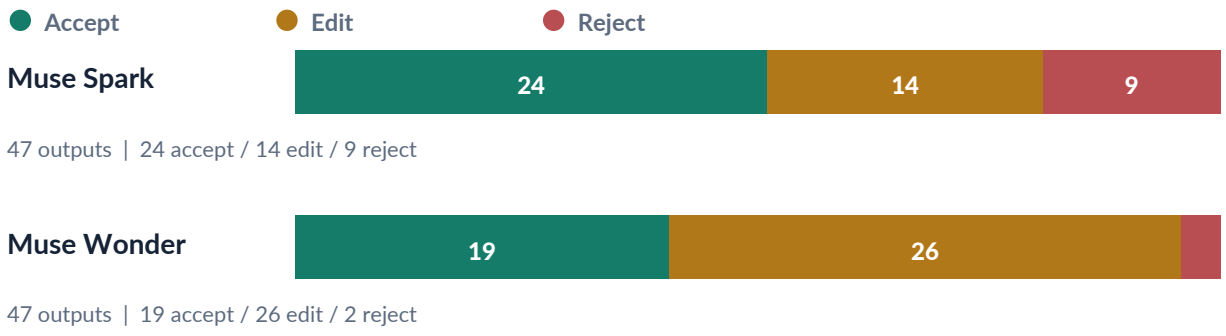
**95**
MODEL CALLS

**66,558**
TOTAL TOKENS

**$0.6567**
EST. MODEL COST

**94**
REVIEWED OUTPUTS

## OWNER VERDICTS - COMPARING ACTUAL USEFULNESS




47 outputs | 24 accept / 14 edit / 9 reject

47 outputs | 19 accept / 26 edit / 2 reject

**Combined:** 43 accept (45.7%) · 40 edit (42.6%) · 11 reject (11.7%)

### Agreement is not accuracy

Agreement only asks whether Spark and Wonder match. Owner verdicts ask whether an output should be used, corrected, or rejected.

The 94 reviewed outputs come from 47 paired evaluations over 19 distinct mail and job examples, tested across versions and settings. A separate routing call brings the total to 95.

**Cost is an estimate, not a bill: the public configuration uses placeholder Spark prices.**

05 / WHY THE OWNER MATTERS

# AI can look right and still be wrong.

A documented model mistake shows why human approval matters.

REAL BENCHMARK ARTIFACT

## $2,000/year

A model explanation referenced this incorrect salary floor while evaluating job fit. The actual standing salary threshold was much higher.

### WHAT YOUR PROCESS DOES INSTEAD

01

**PRESERVE**

Keep the original model output unchanged.

02

**DISCLOSE**

Explain the mistake in public limitations.

03

**DECIDE**

Let the owner accept, edit, or reject the output.

### WHAT THE PUBLIC PROJECT PROVES - AND WHAT IT DOES NOT

PUBLICLY INSPECTABLE

Python router; mail and job benchmark lanes; logs; verdict labels; a dashboard that calculates totals from recorded data.

PUBLIC EVIDENCE BOUNDARY

The private production fleet and approval gates are not exposed in this repo. The day brief shown is a **real fleet-generated output**; its appearance on the physical e-ink panel has not been publicly demonstrated.

06 / HOW TO EXPLAIN IT

# A plain-language explanation

Companion wording, not a replacement for your approved video script v3.

> Suite Fleet is my **personal AI workforce**. I direct it, and specialized AI agents handle everyday tasks like sorting email, scoring job opportunities, reviewing notes, and evaluating deals.
>
> For this hackathon, I integrated **Meta's Muse Spark and Muse Wonder** into workflows I already use. I built a model-routing and measurement layer that records speed, estimated cost, and results on the same real tasks.
>
> The important part is **accountability**. I can compare outputs, see mistakes, and judge what should be accepted, edited, or rejected. AI handles much of the preparation and technical work, but **I keep control over consequential decisions.**
>
> Suite Fleet is not about removing the human. It is about giving one person the capabilities of a much larger team while keeping that person in charge.

WHAT MAKES THE IDEA DISTINCTIVE

**The innovation is not inventing multi-agent AI.**

It is making an existing personal workflow measurable, auditable, and answerable to its owner.

**Bottom line:** You are not building an AI that replaces you. You are directing an AI organization that works for you.

PROJECT LINKS

Live demo  Results  GitHub